

Concurrency Patterns and Mutexes

DD1396

Ric Glassey

glassey@kth.se

Concepts

Concurrency Patterns

Explicit Locking with Mutexes

Concurrency Patterns

Patterns

Patterns are **reusable solutions**

- Design patterns
- Architectural patterns
- Antipatterns

Concurrency patterns


- **Fan-in**
- **Timeout**
- **Replication**

Google Search in Go

Draw: Abstract Google

Queries

Google



Fast ✓
Available ✓
Redundancy ✓

Search verticals

```
type Result string

type Search func(query string) Result

var (
    // search "verticals" modelled as functions
    Web = fakeSearch("web")
    Image = fakeSearch("image")
    Video = fakeSearch("video")
)

func fakeSearch(kind string) Search {
    // factory function to create a search vertical function
    return func(query string) Result {
        // fake search that just waits a few milliseconds
        time.Sleep(time.Duration(rand.Intn(100)) * time.Millisecond)
        return Result(fmt.Sprintf("%s result for %q\n", kind, query))
    }
}
```

Execute a search

```
func Google(query string) ([]Result) {  
    // Search for Web, Image and Video results  
    results = append(results, Web(query))  
    results = append(results, Image(query))  
    results = append(results, Video(query))  
    return  
}  
  
func main() {  
    // Time the search  
    rand.Seed(time.Now().UnixNano())  
    start := time.Now()  
    results := Google("golang")  
    elapsed := time.Since(start)  
    fmt.Println(results)  
    fmt.Println(elapsed)  
}
```

Draw: Sequential Search

Sequential



Demo: Sequential Search

Sequential slowdown

Each search vertical is called in a **linear sequence**

- May have different **performance**
- May have different **running conditions**


Opportunity for **concurrency**

- Search verticals are waiting to be called
- Runtime will be the sum of all three systems

How can this be improved (using Go)?

Fan-in Pattern

Fan-In



Fan in pattern




Fan-in

```
func Google(query string) ([]Result) {  
    // Use a channel to collect results  
    comm := make(chan Result)  
  
    // create three search threads using a fan-in pattern  
    go func() { comm <- Web(query) } ()  
    go func() { comm <- Image(query) } ()  
    go func() { comm <- Video(query) } ()  
  
    // collect results  
    for i := 0; i < 3; i++ {  
        result := <- comm  
        results = append(results, result)  
    }  
    return  
}
```

Demo: Fan-in

Time-out Pattern

Timeout



Time-out

```
func Google(query string) ([]Result) {
    comm := make(chan Result)


    // create three search threads using a fan-in pattern
    go func() { comm <- Web(query) } ()
    go func() { comm <- Image(query) } ()
    go func() { comm <- Video(query) } ()

    // collect results; but do not wait on slow services
    timeout := time.After(70 * time.Millisecond)
    for i := 0; i < 3; i++ {
        select {
        case result := <- comm:
            results = append(results, result)
        case <- timeout:
            fmt.Println("timed out")
            return
        }
    }
    return
}
```


Demo: Time-out Pattern

Replication Pattern

Replication



Replication

```
func Google(query string) ([]Result) {
    comm := make(chan Result)

    // create three search threads using a fan-in pattern
    go func() { comm <- First(query, Web1, Web2) } ()
    go func() { comm <- First(query, Image1, Image2) } ()
    go func() { comm <- First(query, Video1, Video2) } ()
    return
}

func First(query string, replicas ...Search) Result {
    // launch replicas and return fastest response
    c := make(chan Result)
    searchReplica := func(i int) { c <- replicas[i](query) }
    for i := range replicas {
        go searchReplica(i)
    }
    return <-c
}
```

Demo: Replication

Complex system; Low cost

Modern scalable systems need to be

- Fast
- Available
- Redundant

Go provides the concurrency primitives

Concurrency pattern provides a solution

Locking by Mutual Exclusion

Simple is better

Race conditions are a problem of concurrent systems

Go can share variables over channels as one solution

The classic solution is **mutual exclusion**

- Lock a variable
- Only one thread can have write access
- “No two threads are in their **critical section** at the same time”
- i.e. reading, updating and writing a shared variable

Using a mutex

Mutex is an element within a program that can **control access to shared data** by **locking** and **unlocking**

During critical section (e.g. update):

- Lock is acquired by thread
- Work is done, then
- Lock released for other waiting threads

My First Bank App

```
type account struct {  
    amount    float64  
}  
  
func (acc *account) Deposit(sum float64) {  
    acc.amount += sum  
}  
  
func (acc *account) Withdraw(sum float64) {  
    acc.amount -= sum  
}  
  
func (acc *account) Balance() string {  
    return strconv.FormatFloat(acc.amount, 'f', 2, 64) + " Kr"  
}
```

Simulation setup

```
func main() {  
  
    var joint_account account  
    joint_account.Deposit(1000.00)  
  
    // stop main from quitting before threads  
    wg := new(sync.WaitGroup)  
    wg.Add(2)  
  
    // Concurrent access of joint bank account (what could possibly go wrong?)  
    go func () { // Person One  
        joint_account.Deposit(50.00)  
        joint_account.Deposit(50.00)  
        joint_account.Withdraw(200.00)  
        joint_account.Deposit(50.00)  
        wg.Done()  
    }()  
  
    go func () { // Person Two  
        joint_account.Deposit(50.00)  
        joint_account.Deposit(50.00)  
        joint_account.Withdraw(200.00)  
        joint_account.Deposit(50.00)  
        wg.Done()  
    }()  
  
    // ensure both customers have finished  
    wg.Wait()  
    fmt.Println(joint_account.Balance())  
}
```

Demo: My First Bank App

Detecting Race Conditions

Runtime support

Go runtime can be asked to check for race conditions

```
$ go run -race myapp.go
```

Why not the default?

- “The cost of race detection varies by program, but for a typical program, memory usage may increase by 5-10x and execution time by 2-20x” https://golang.org/doc/articles/race_detector#Runtime_Overheads

Mutex in Go

Import the **sync** package

Create a struct to hold the mutex and variable

Mutex has two methods: **Lock** and **Unlock**

Goroutines wait before accessing the variable until they have the lock

My Better Bank App

```
import (  
    "fmt"  
    "strconv"  
    "sync"  
)  
  
type account struct {  
    mu      sync.Mutex  
    amount float64  
}
```

Locking critical sections

```
func (acc *account) Deposit(sum float64) {  
    acc.mu.Lock()  
    acc.amount += sum  
    acc.mu.Unlock()  
}
```

```
func (acc *account) Withdraw(sum float64) {  
    acc.mu.Lock()  
    acc.amount -= sum  
    acc.mu.Unlock()  
}
```

Demo: My Better Bank App

Concept Review

Concept review

Go supports complex concurrent systems with compact code

Common concurrency patterns can be elegantly implemented via

- Goroutines
- Channels
- Select
- Timeouts

Yet...there are still practical situations where channels are not suited, and a traditional locking by mutual exclusion is preferred

Recommended Reading

Any Rob Pike talk on Go & Concurrency :-)

- <https://www.youtube.com/watch?v=f6kdp27TYZs>

Fundamentals of concurrent programming

- Mutual Exclusion lock (mutex)
- <https://yourbasic.org/golang/mutex-explained/>